



Wav2Letter - Implementing an open source E2E ASR in PyTorch

Assaf Mushkin



Everything you need to know

- Wav2Letter is an easy to understand, accurate model for ASR
- There are many improvements on it
- I've made a simple, minimalist PyTorch implementation available online




Agenda

- CTC and Wav2Letter
- Spinoffs and improvements
- Lessons from working with open source projects
- Live Demonstration!


Wav2Letter (2016 - FAIR)

- Input spectrograms, Output is CTC (probability of each letter per frame)



Wav2Letter (2016 - FAIR)


- Input spectrograms, Output is CTC (probability of each letter per frame)
- Fully convolutional, just a long stack of convolutions



Wav2Letter



A	0.9	0.1	0.3	0
B	0.1	0.7	0.1	0.1
C	0	0.2	0.6	0.9

Text



Wav2Letter (2016 - FAIR)

- Input spectrograms, Output is CTC (probability of each letter per frame)
- Fully convolutional, just a long stack of convolutions
- MFCC features (spectrogram - only slightly worse)





CTC - “Connectionist Temporal Classification”

- For each frame / window, predict probability (or score) of each letter.
 - Add unique “blank” character, _
- When interpreting, drop consecutive identical letters, then drop blanks.

Frame	1	2	3	4	5
A	0.6	0.3
B	0.3	0.2
blank	0.1	0.5

_ABA -> ABA
AA_BA -> ABA
AB_BA -> ABBA
AAA_BB_A_ -> ABA



CTC example

CTC : HH_E_LLLLL_LL_OO HH_O_WWW ARRRR____EE YOUUU

OUTPUT: HELLO HOW ARE YOU



CTC - cont.

- CTC can be interpreted “greedily” - pick letter for each frame independently, then run reduction

A	0.9	0.1	0.3	0
B	0.1	0.7	0.1	0.1
_	0	0.2	0.6	0.9



CTC - cont.

- CTC can be interpreted “greedily” - pick letter for each frame independently, then run reduction
- Can also be decoded exhaustive - find output string with highest total score, considering all possible reductions



CTC - cont.

- CTC can be interpreted “greedily” - pick letter for each frame independently, then run reduction
- Can also be decoded exhaustive - find output string with highest total score, considering all possible reductions
- Greedy results aren’t the best, but are much faster to compute



CTC - cont.

- CTC can be interpreted “greedily” - pick letter for each frame independently, then run reduction
- Can also be decoded exhaustive - find output string with highest total score, considering all possible reductions
- Greedy results aren’t the best, but are much faster to compute
- Practically, we use beam search, weighed with a LM
 - This adds a bunch more hyperparameters to the process, can be found independently of the acoustic model



ASG - CTC but with a twist

- Used in Wav2letter (2016)





ASG - CTC but with a twist

- Used in Wav2letter (2016)
- TL;DR: Rarely used in future works, use CTC instead.



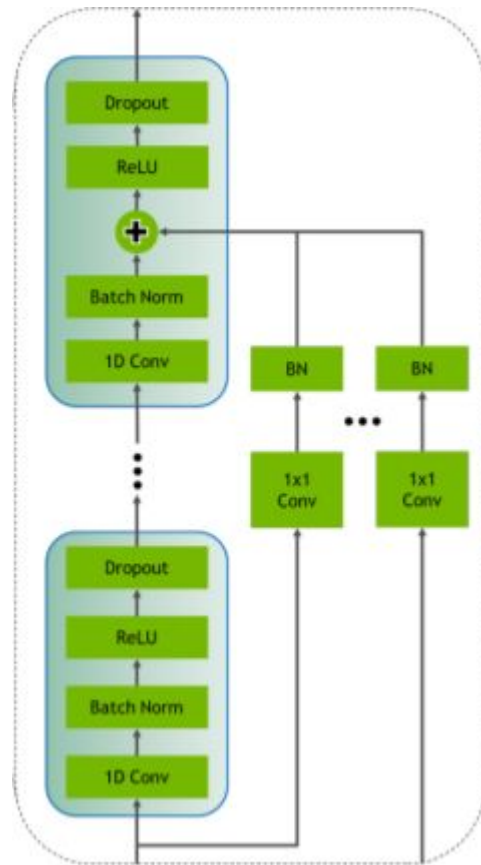
ASG - CTC but with a twist

- Used in Wav2letter (2016)
- TL;DR: Rarely used in future works, use CTC instead.
- No usage of blank char, instead repeat last character and use 2 and 3 for repetition (“**bo2k**” instead of “**book**”)
 - This defines a different “reduce” function, and a different simpler decoding algorithm
- Global normalization instead of per-frame normalization
 - This isn’t connected to ASG necessarily, but was in the same paper.



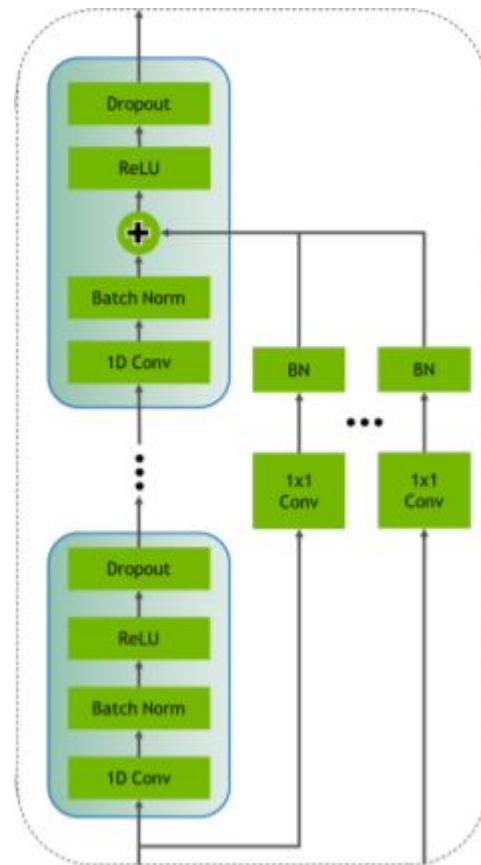
Jasper (2019) - Bigger! Better!

- Pretty much Wav2Letter, but with residual connections
- “JasperBlock”:
 - **R** times: [Conv1d,BN, Relu,Dropout]
 - + residual connection
- Jasper architecture is **B** repetitions of JasperBlock



Jasper (2019) - Bigger! Better!

- Pretty much Wav2Letter, but with residual connections
- “JasperBlock”:
 - R times: [Conv1d, BN, Relu, Dropout]
 - + residual connection
- Jasper architecture is B repetitions of JasperBlock
- From the paper, $B=10$ and $R=3$ (LibriSpeech) or $R=5$ (WSJ)
- $3 \times 10 = 30$ convolutions, compared to 12 in Wav2letter



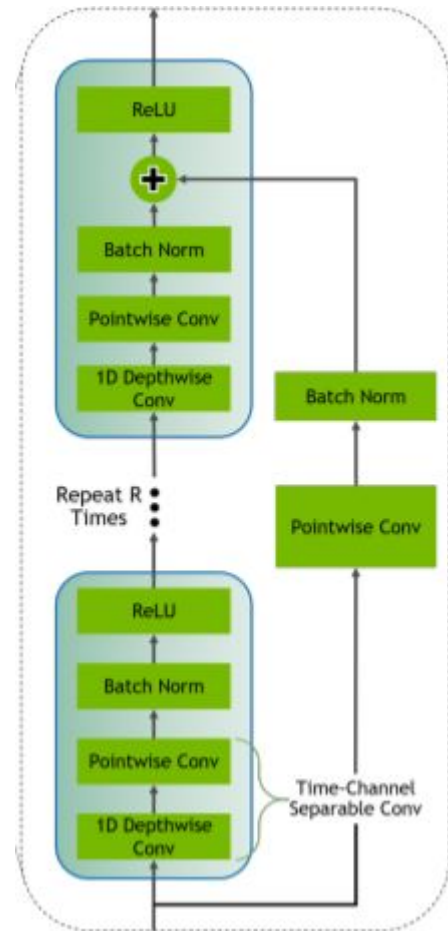


Novograd - Jasper's Optimizer

- It's Adam, but momentum is calculated per layer instead of per weight
- Sometimes better end results, but mostly just speedup for ASR task (anecdotally, 30% less epochs to reach same loss)
- Current rumour: No speedup in other tasks, compared to Adam
- Not implemented in PyTorch yet, but is open-source and available online

QuartzNet (2019 - 1 month later)

- Jasper + 1D Time-Depth-Separable convolutions *




* First use of TDS was in FAIR in a different article

QuartzNet (2019 - 1 month later)

- Jasper + 1D Time-Depth-Separable convolutions *
- Instead of passing each convolution a matrix of size [window X channels], break into two convolutions.
- First convolve over time (treating each channel the same), then convolve over channels with a window of 1. Called depthwise and pointwise convolutions

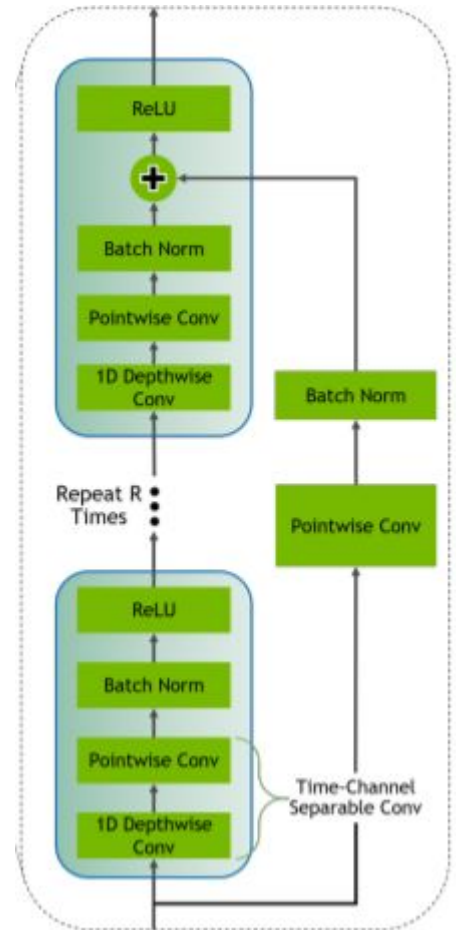
* First use of TDS was in FAIR in a different article



QuartzNet (2019 - 1 month later)

- Jasper + 1D Time-Depth-Separable convolutions *
- Instead of passing each convolution a matrix of size [window X channels], break into two convolutions.
- First convolve over time (treating each channel the same), then convolve over channels with a window of 1. Called depthwise and pointwise convolutions
- Drastically decrease the number of parameters (20 mil. Instead of 100)

* First use of TDS was in FAIR in a different article





Wav2letter improvements

- Shorter future context, speedups - better for online decoding ([link](#))
- Semi-supervised training ([link](#)) - also included ResNet style acoustic models
- Lexicon free decoding ([link](#)) - better performance on OOV





Sidenote - integrating open source projects

- Actually not that hard



Sidenote - integrating open source projects

- Actually not that hard
- Projects are not libraries, different mindset required



Sidenote - integrating open source projects

- Actually not that hard
- Projects are not libraries, different mindset required
- It will not work out-of-the-box, you will have to read and edit the code



Sidenote - integrating open source projects

- Actually not that hard
- Projects are not libraries, different mindset required
- It will not work out-of-the-box, you will have to read and edit the code
- They made mistakes, bad assumptions, or are out-of-date



Sidenote - integrating open source projects

- Actually not that hard
- Projects are not libraries, different mindset required
- It will not work out-of-the-box, you will have to read and edit the code
- They made mistakes, bad assumptions, or are out-of-date
- The price we pay for cutting-edge developments



Show, don't tell

- Demonstration training for a very small dataset
 - Tensorboard is useful for visualizing training process
-
- Repository available at https://github.com/assafmu/wav2letter_pytorch
 - PyTorch implementation, intentionally minimalist